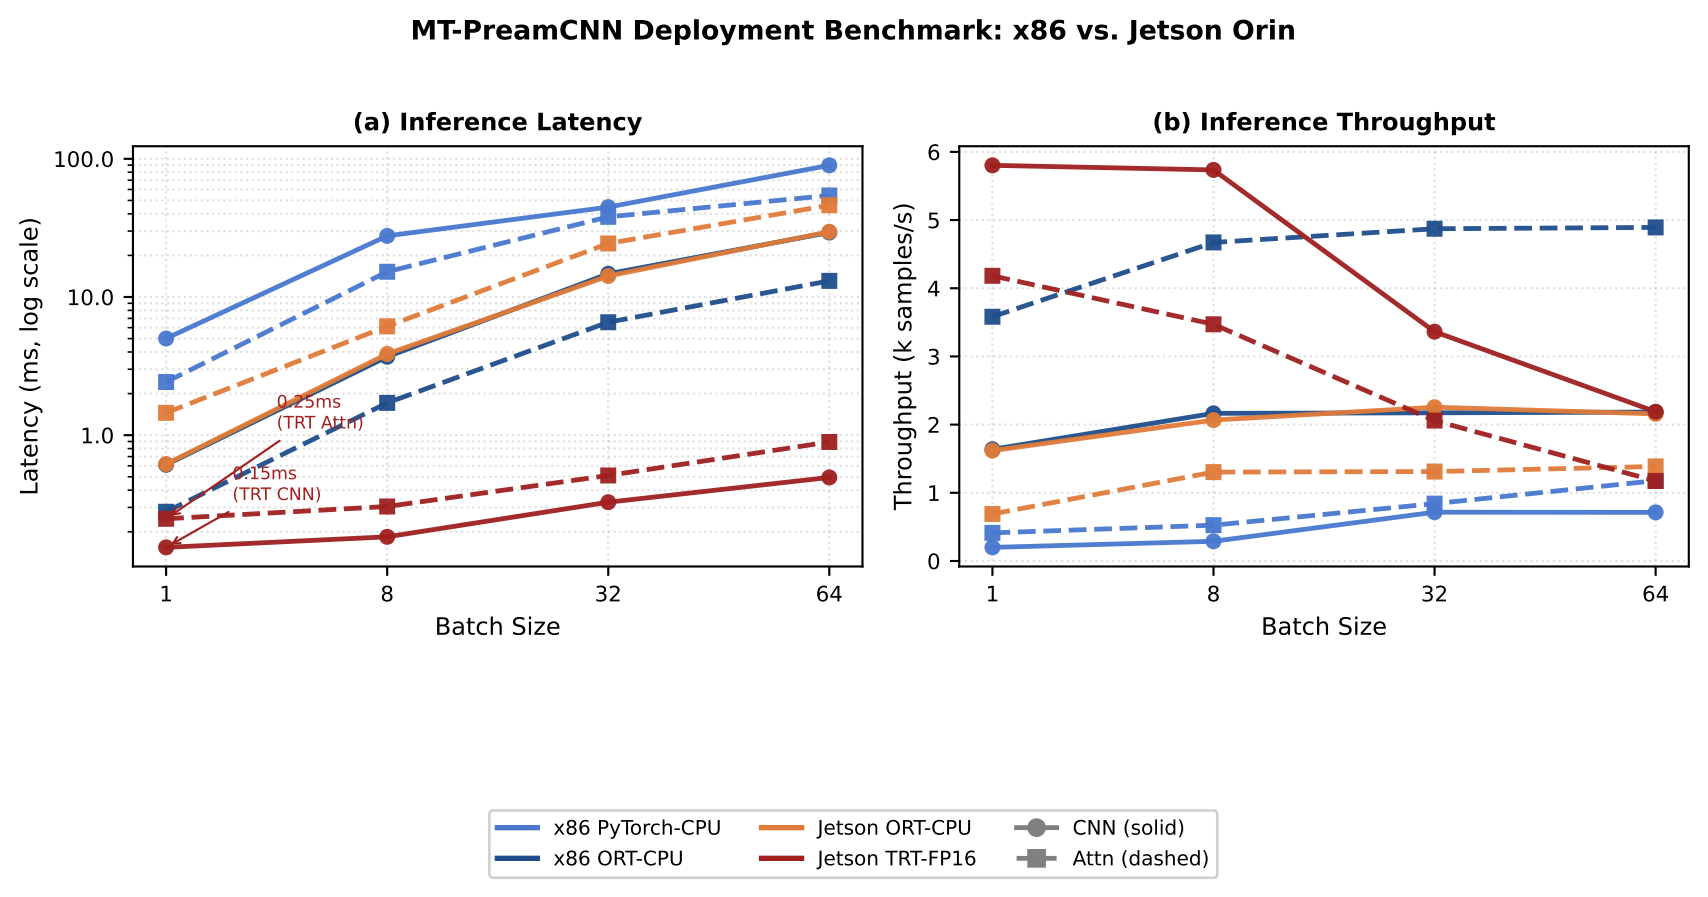
MT-PreamCNN Deployment Benchmark: x86 vs. Jetson Orin
(a) Inference Latency
Latency (ms, log scale)
100.0
10.0
1.0
0.25ms
(TRT Attn)
0.15ms
(TRT CNN)
1
8
32
64
Batch Size
(b) Inference Throughput
Throughput (k samples/s)
0
1
2
3
4
5
6
1
8
32
64
Batch Size
x86 PyTorch-CPU
x86 ORT-CPU
Jetson ORT-CPU
Jetson TRT-FP16
CNN (solid)
Attn (dashed)