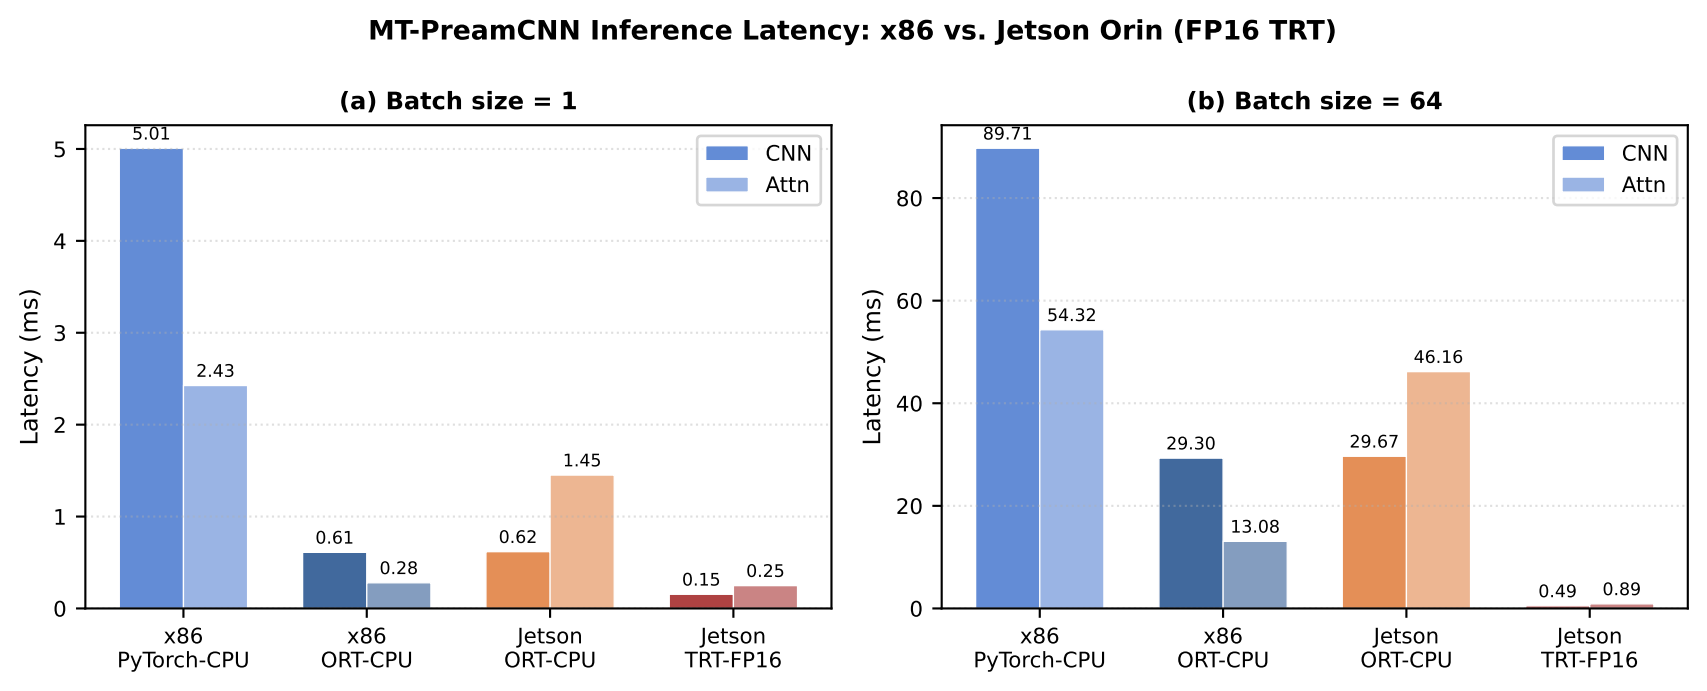
MT-PreamCNN Inference Latency: x86 vs. Jetson Orin (FP16 TRT)
(a) Batch size = 1
(b) Batch size = 64
Latency (ms)
CNN
Attn
5.01
2.43
0.61
0.28
0.62
1.45
0.15
0.25
89.71
54.32
29.30
13.08
29.67
46.16
0.49
0.89
x86 PyTorch-CPU
x86 ORT-CPU
Jetson ORT-CPU
Jetson TRT-FP16